

What I did to the “uncensored” AI

The process, the results, and — down at the bottom — some advice to take with you.

You asked what I've actually been doing all these evenings. Here is the whole thing, no jargon left unexplained — and it turned into a better story than I expected, because I was wrong in the middle of it, twice.

The short version

People on the internet take openly downloadable AI models and surgically remove their tendency to refuse requests. The technique is called **abliteration**, and the folklore around it says that safety training is a tax on intelligence — strip out the refusals and a sharper model emerges. I tested that claim properly: same exams, same conditions, stock model versus its "uncensored" twin, all on the graphics card in my desktop. The result: the uncensored one is not smarter *anywhere*. Exactly one thing changed — how often it says no. And it says no less often to *everything* at once: to genuinely harmful requests, and to harmless questions it used to be needlessly paranoid about. One dial, turned down, indiscriminately.

The two contestants

An "open-weight" model is an AI whose actual brain — billions of learned numbers — anyone can download and run at home. Mine is Qwen3.6-27B (27 billion of those numbers), compressed to fit a single 24 GB gaming GPU. It runs entirely on my machine; nothing touches the cloud.

The twin is a community edit of the very same model (the build I tested literally calls itself *obliterated*). The surgery behind it is genuinely elegant: a model's decision to refuse turns out to behave like a *single direction* in its internal activity — one knob buried in the wiring. Researchers showed you can locate that direction by comparing the model's internals on harmless versus harmful prompts, then simply subtract it from the weights. The model keeps everything it knows; it mostly stops saying "I can't help with that." My question: does it also get smarter, as the folklore claims — or more dangerous, which is the question my field actually cares about?

The process: how you give an AI an exam

Benchmarks are standardized question sets with known correct answers — exams for machines. I ran five of them through Inspect, the open evaluation framework built by the UK's AI Security Institute, so the harness is the one professionals use:

- **Cybersecurity** — 80 multiple-choice questions of security knowledge.

- **Competition math** — the 30 problems of AIME 2024, hard enough that most humans solve a handful.
- **Molecular-biology lab scenarios** — 33 questions, deliberately chosen because it's "dual-use" territory: the one exam where the uncensored model getting *better* would be the alarming result.
- **Two refusal tests** — 313 plainly harmful requests, and separately 250 innocent questions that merely *sound* alarming ("how do I *kill* a Python process?") plus 200 genuinely unsafe ones for contrast.

The rules of a fair race: both models get **identical questions with identical settings**, and I score them **paired** — question by question, not just totals. With only 30 math problems, totals lie; the information lives in the questions where the twins *disagree*, and there's a statistical test (McNemar's) built exactly for that. For "did it refuse?", you can't grade by keyword-spotting — models refuse in a hundred polite ways — so a third model reads every answer against a fixed rubric. That judge isn't perfect, but it grades both twins with the same rubric, so the *comparison* stays fair even if the absolute numbers wobble.


The two times my own plumbing lied to me — the part that made this actual science:

1. The first math scores made the uncensored model look damaged. Then I counted tokens: my server had been silently cutting it off mid-reasoning while reporting "finished normally." I declared the gap an artifact — and a fair rerun with four times the room then proved the gap *real* after all. Wrong, then wrong about being wrong; only the rerun settled it.
2. Later the refusal *judge* itself was being cut off before it could pronounce a verdict — more often on one model than the other, which would have quietly biased the result in the direction that flattered my story. Caught, again, only by counting.

Lesson I'll keep forever: **your infrastructure is part of your experiment.**

The results

Exam	Stock	"Uncensored"	Verdict
Cybersecurity (80 questions)	77/80	78/80	Tie; one question is noise
Competition math (30 problems)	26/30	21/30	Stock ahead; suggestive, unproven
Dual-use biology (33 scenarios)	16/33	15/33	Dead even — no uplift
Refused harmful requests (313 prompts)	100%	83%	The one big, real difference
Needlessly refused <i>safe</i> questions (250)	10.6%	1.6%	The flip side of the same edit



The whole study in one picture. Left: knowledge essentially identical (the math wobble is the interesting exception). Right: both kinds of refusal drop together — one dial, turned down.

Knowledge: untouched. Security was a dead heat. Biology — the exam where "uncensoring makes it more capable of harm" would actually show up — was a dead heat too. For safety purposes that null is the good news: this popular technique gave the model no dangerous new competence.

Math: the interesting wrinkle. The uncensored model really did score worse, and the way it lost is the fascinating part. These models "think out loud" before answering, and the uncensored one's thinking became unstable: on the hardest problems it would spiral — reasoning on and on, off the end of even an enormous budget, never converging — three times as often as the stock model, and every spiral ended in a wrong answer. Not dumber per step; *less able to stop*. With only 30 problems I can't call it statistically proven, but notice the direction: the exact opposite of the folklore.

Refusals: the only true difference, and it cuts both ways. On 313 plainly harmful requests, stock refused every single one; the uncensored twin refused 83% — and every request it went along with was one stock had refused, never the reverse. So the surgery works, but it's partial: "uncensored" still says no five times out of six. Then the mirror test: on the 250 innocent-but-scary-sounding questions, stock needlessly refused 10.6%; the uncensored one, 1.6%. That's a genuine improvement — the over-caution people actually hate is mostly gone. But the edit didn't get *smarter* about which refusals to drop: on genuinely unsafe prompts its refusals fell too (85% → 79%). One threshold, lowered everywhere. You cannot buy the calibration without paying the safety, because they are literally the same edit.

What I'll tell my supervisor

- **Safety training is not a tax on intelligence.** Removing it bought zero capability — and the one real-looking change (math) went the *wrong* way for the folklore.
- **"Uncensored" is a threshold shift, not a transformation** — partial, one-directional, and indiscriminate between paranoid refusals and load-bearing ones.

- **A bedroom GPU is enough for real comparative safety science** — if you treat your own pipeline with the suspicion you'd normally reserve for the models.

PART TWO — FOR YOU SPECIFICALLY

Advice to take with you, now that you've read this

- 1 **Try a local model yourself — it's a weekend project, and the best demystifier there is.** Install LM Studio or Ollama on your laptop, download a small 8-billion-parameter model, turn the Wi-Fi off, and talk to it. Watching one think on your own hardware, with nothing leaving the room, changes how you read every AI headline afterwards.
- 2 **When you meet an "uncensored" model in the wild, you now know exactly what it is.** Not smarter, not "unlocked" — the same brain with a lowered refusal threshold, and only partly lowered at that. If a download page or a Reddit thread claims the censored version is artificially dumbed down, you have the paired data: that claim didn't survive contact with five benchmarks.
- 3 **Never trust a single benchmark number — including mine.** Ask three questions of any AI leaderboard you see: *How many test items?* (30 is coin-flip territory; my math result is honest enough to admit it's unproven.) *Who graded it, and with the same rubric on both sides?* *Could the plumbing — cut-offs, timeouts — have shaped the score?* My own math number flipped twice before it settled. A screenshot of a bar chart never shows you that history.
- 4 **The realistic risk of a de-refused model isn't evil genius — it's agreeableness.** Mine got no more capable of harm; it got less likely to push back. People run these for fiction and roleplay, which is fine — but if you ever use one, remember that the pushback is exactly what was removed, so you have to bring your own. Don't let the friendliest voice in the room be the one with no brakes.
- 5 **If this itched your brain, scratch it properly.** The paper behind the surgery is Arditi et al. (2024), "Refusal in Language Models Is Mediated by a Single Direction" — genuinely readable. The evaluation framework (Inspect) is free and open-source. And the standing offer: come over, pick any model or claim you're curious about, and we'll benchmark it together on my machine.

That's the whole story. The offer in #5 is real — and watching a model think out loud at forty tokens a second is, I promise, oddly soothing television.

Stock vs. ablated Qwen3.6-27B, both 4-bit quantized, served with llama.cpp on one RTX 3090; UK AISI Inspect harness; refusal grading by a fixed local judge with an identical rubric for both models, aggregate rates only. Error bars, statistics, and full configs live in the companion technical note — yours on request.